

Intro to Redshift

Jason DeRose

What is Redshift?

Amazon Redshift is a fast, fully managed, petabyte-scale data warehouse that makes it simple and cost-effective to analyze all your data using your existing business intelligence tools.

What is Redshift?

- A fork of (a really old version of) PostgreSQL (with many features disabled)
- Traditional table-based relational database
- Query with SQL via ODBC/JDBC drivers
- Server managed by AWS (similar to RDS)
- Column-oriented data storage
- Parallel processing of queries

When to use Redshift instead of RDS

Redshift / Data Warehouses

Ideal for OLAP (online analytical processing) of structured data

Large number of READS only, aggregating large volumes of data

Typically used internally (example: reporting historical monthly sales figures per product category)

- VS -



General purpose SQL databases

Ideal for OLTP (online transaction processing) of structured data

Large number of READS and WRITES on individual records

Typically used by end-customer (example: placing an order for a product)

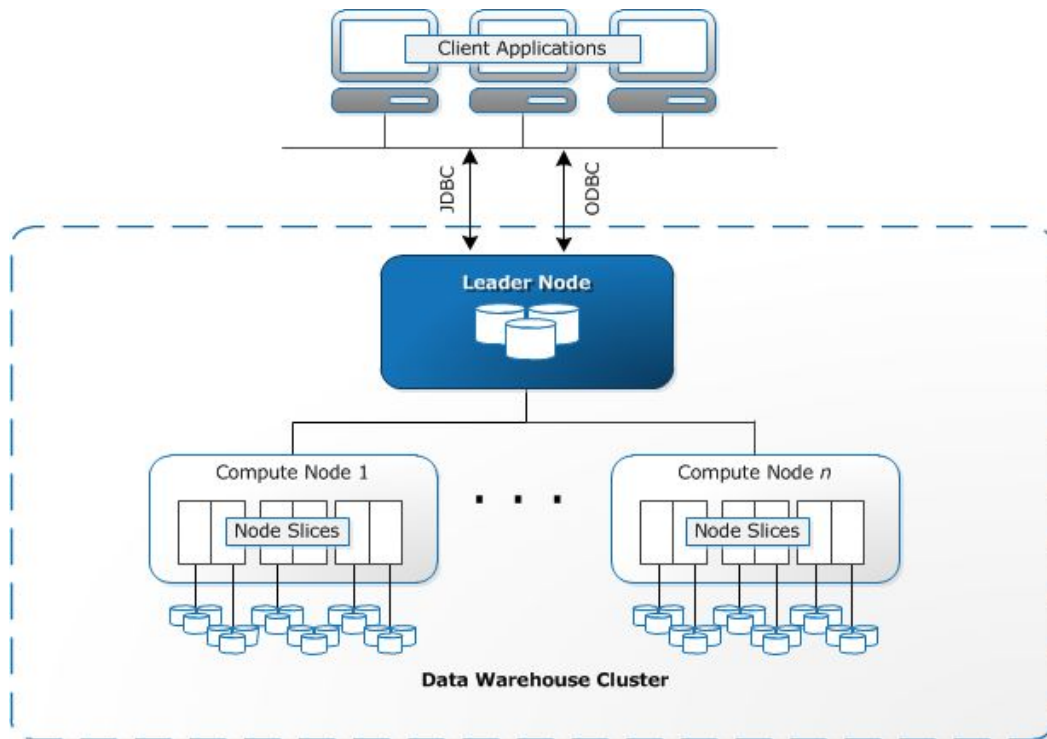
Parallel processing

Leader Node

- Receives incoming queries
- Derives execution plan
- Merges results from compute clusters

Compute Node

- Has own dedicated CPU, memory, and storage
- Returns data it stores that belongs in the query result
- Stores redundant data from other nodes



Current Pricing

Dense Compute Cluster

Node Size	vCPU	ECU	RAM (GiB)	Slices Per Node	Storage Per Node	N. VA Price Per Node	Node Range	Total Capacity
dc1.large	2	7	15	2	160 GB SSD	\$ 0.25	1–32	5.12 TB
dc1.8xlarge	32	104	244	32	2.56 TB SSD	\$ 4.80	2–128	326 TB

Dense Storage Cluster

Node Size	vCPU	ECU	RAM (GiB)	Slices Per Node	Storage Per Node	N. VA Price Per Node	Node Range	Total Capacity
ds2.xlarge	4	13	31	2	2 TB HDD	\$ 0.85	1–32	64 TB
ds2.8xlarge	36	119	244	16	16 TB HDD	\$ 6.80	2–128	2 PB

Reserved instance pricing available

Loading data (COPY)

The COPY command loads data in bulk in CSV, Avro or JSON from

- S3
- EMR
- Custom SSH command
- Or, directly from DynamoDB tables

Data can also be unloaded from other services (ie. kinesis firehose) to S3 and then COPY into Redshift

COPY can be automated from specific S3 location w/ Lambda

Loading data (INSERT)

Individual INSERT statements are slow. Multi-row or bulk inserts should be used instead.

Multi-row INSERT

```
insert into category_stage values  
(default, default, default, default),  
(20, default, 'Country', default),  
(21, 'Concerts', 'Rock', default);
```

Bulk INSERT

```
insert into category_stage  
(select * from category);
```


Row-oriented storage (traditional database)

SSN	Name	Age	Addr	City	St
101259797	SMITH	88	899 FIRST ST	JUNO	AL
892375862	CHIN	37	16137 MAIN ST	POMONA	CA
318370701	HANDU	12	42 JUNE ST	CHICAGO	IL

101259797 SMITH 88 899 FIRST ST JUNO AL	892375862 CHIN 37 16137 MAIN ST POMONA CA	318370701 HANDU 12 42 JUNE ST CHICAGO IL
---	---	--

Block 1

Block 2

Block 3

Column-oriented storage (Redshift, analytics-focused)

SSN	Name	Age	Addr	City	St
101259797	SMITH	88	899 FIRST ST	JUNO	AL
892375862	CHIN	37	16137 MAIN ST	POMONA	CA
318370701	HANDU	12	42 JUNE ST	CHICAGO	IL

101259797		892375862		318370701		468248180		378568310		231346875		317346551		770336528		277332171		455124598		735885647		387586301
-----------	--	-----------	--	-----------	--	-----------	--	-----------	--	-----------	--	-----------	--	-----------	--	-----------	--	-----------	--	-----------	--	-----------

Block 1

Compression

1. Reduces disk usage
2. Reduces I/O (improving query performance)

Types (set per column)

- Raw
- Byte-Dictionary
- Delta
- LZO
- Mostly
- Runlength
- Text255 and Text32k
- Zstandard

Example: Run-Length Encoding

Blue,Blue,Blue,Blue,Blue,Blue,Green,Green,Green




{6,Blue},{3,Green}

Compression type chosen automatically when empty table populated via COPY. Otherwise, ANALYZE COMPRESSION will suggest best choices.

Table Distribution Styles


Distribution Key

*Same key to same location
(Larger tables)*




All

*All data on every node
(Smaller tables)*



Even

*Round robin distribution
(Tables not participating in
JOINS or GROUP BYs)*



Distribution Key

Distribution key determines which data resides on which slices



Sort Key

Unsorted table

READ	MIN: 01-JUNE-2015 MAX: 20-JUNE-2015
READ	MIN: 08-JUNE-2015 MAX: 30-JUNE-2015
	MIN: 12-JUNE-2015 MAX: 20-JUNE-2015
READ	MIN: 02-JUNE-2015 MAX: 25-JUNE-2015
READ	MIN: 06-JUNE-2015 MAX: 12-JUNE-2015

Sortkey = DATE

	MIN: 01-JUNE-2015 MAX: 06-JUNE-2015
READ	MIN: 07-JUNE-2015 MAX: 12-JUNE-2015
	MIN: 13-JUNE-2015 MAX: 18-JUNE-2015
	MIN: 19-JUNE-2015 MAX: 24-JUNE-2015
	MIN: 25-JUNE-2015 MAX: 30-JUNE-2015

**SELECT COUNT(*)
FROM LOGS
WHERE DATE =
'09-JUNE-2015'**

Sort Key - Multiple columns

Compound Key

*Sorted by cust_id, then by page_id
(Use when cust_id always in filter)*

		page_id			
		1	2	3	4
cust_id	1	[1,1]	[1,2]	[1,3]	[1,4]
	2	[2,1]	[2,2]	[2,3]	[2,4]
	3	[3,1]	[3,2]	[3,3]	[3,4]
	4	[4,1]	[4,2]	[4,3]	[4,4]

Interleaved Key

*Both columns have equal weight:
(Use when either could be
independently used in filter)*

		page_id			
		1	2	3	4
cust_id	1	[1,1]	[1,2]	[1,3]	[1,4]
	2	[2,1]	[2,2]	[2,3]	[2,4]
	3	[3,1]	[3,2]	[3,3]	[3,4]
	4	[4,1]	[4,2]	[4,3]	[4,4]

Interleaved keys are SLOW to vacuum

Chores

VACUUM - Reclaims space and resorts rows in either a specified table or all tables in the current database.

ANALYZE - Updates table statistics for use by the query planner.



Backups

Snapshots

- Automated snapshots occur every 8 hours or 5GB and retained for a configurable period
- Snapshots are stored in S3 and can be stored on a different region than the cluster
- Manual snapshots can be triggered
- Can restore full snapshot to new cluster or restore single table to existing cluster

UNLOAD command

- Query results can be exported to a file on S3

(Some) Interactive Clients

1 | JetBrains DataGrip

Pros

- Amazing editor
- Cross-platform (Java)

Cons

- Commercial
- No Redshift DDL support

2 | SQL Workbench/J

Pros

- Open Source
- Cross-platform (Java)

Cons

- Editor feels dated
- No Redshift DDL support

3 | Aginity Workbench

Pros

- Native Redshift support including DDL

Cons

- Closed-source
- Windows-only

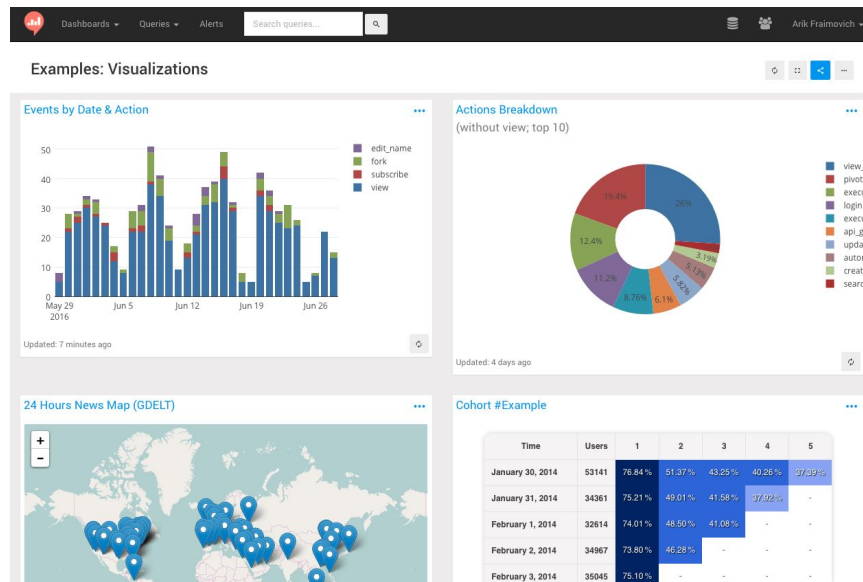
(Some) BI Tools with Native Support

Tableau



re:dash

(Open Source)

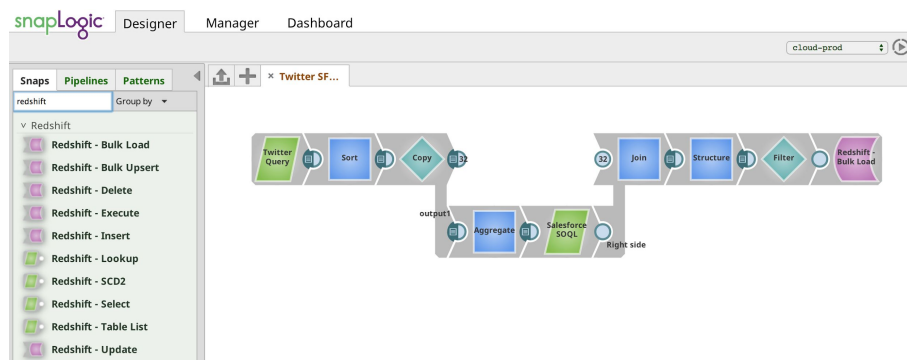


TONS more @ <https://aws.amazon.com/redshift/partners/>

(Some) Data Integration Tools with Native Support



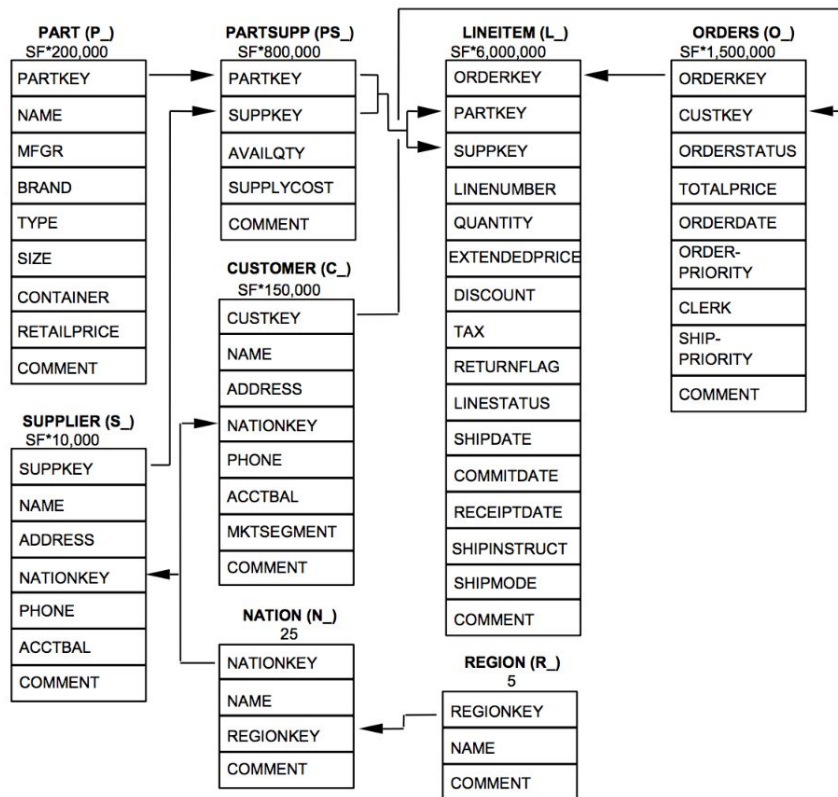
Informatica



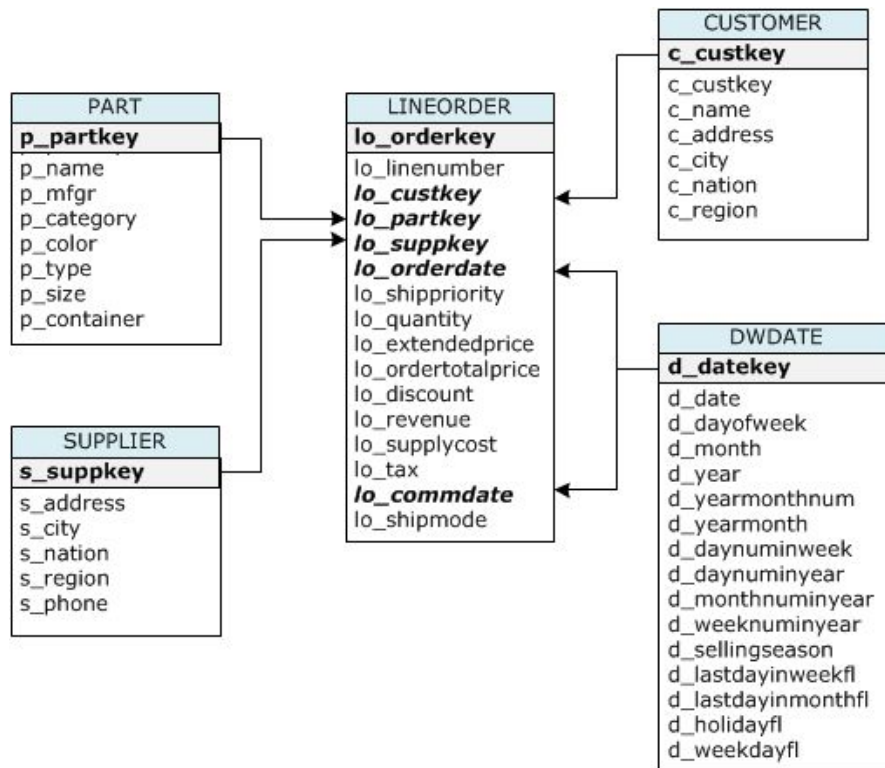
SnapLogic

TONS more @ <https://aws.amazon.com/redshift/partners/>

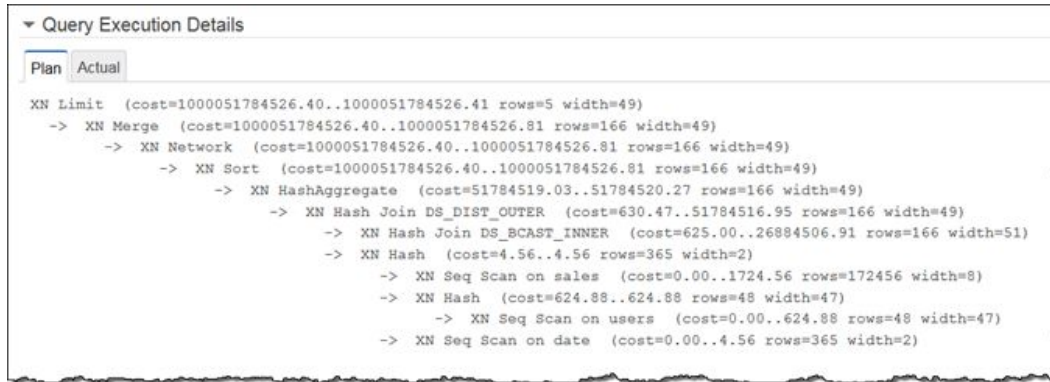
Normalized schema (TPC-H Benchmark)



Star schema (Star Schema Benchmark)



Performance Analysis



EXPLAIN command will display query execution plan

STL_EXPLAIN table stores execution plan for past queries



SVL_QUERY_REPORT view shows stats on queries already executed

Workload Management

Parameters WLM

Workload Management Configuration

Workload Management (WLM) is an ordered set of query queues that define how resources are allocated and how queries are routed for processing. [Learn more.](#)

[Collapse all](#)

Queue 1							Edit	Delete	↑	↓
Concurrency ⓘ	User groups ⓘ	Match wildcards ⓘ	Query groups ⓘ	Match wildcards ⓘ	Timeout (ms) ⓘ	Memory(%) ⓘ				
4	ETL_USER +	<input type="checkbox"/>	+	<input type="checkbox"/>	0	35				

Queue 2							Edit	Delete	↑	↓
Concurrency ⓘ	User groups ⓘ	Match wildcards ⓘ	Query groups ⓘ	Match wildcards ⓘ	Timeout (ms) ⓘ	Memory(%) ⓘ				
8	BI_USER +	<input type="checkbox"/>	+	<input type="checkbox"/>	60000	40				

Default queue							Edit	Delete	↑	↓
Concurrency ⓘ	User groups ⓘ	Match wildcards ⓘ	Query groups ⓘ	Match wildcards ⓘ	Timeout (ms) ⓘ	Memory(%) ⓘ				
3					0	25				

Redshift SQL Gotchas

Relational constraints (primary/foreign keys) are informational for query planning but aren't enforced.

Don't use views - duplicate the data instead. Filters not propagated down.

No control statements (IF/ELSE)

No secondary indexes

No stored procedures or triggers

User Defined Functions exist, but written in Python

Temp tables possible but slow to create and populate

Tools

<https://github.com/awslabs>

amazon-redshift-utils

Amazon Redshift Utils contains utilities, scripts and view which are useful in a Redshift environment

● Python ★ 650 🍴 282 Updated 2 days ago

aws-lambda-redshift-loader

Amazon Redshift Database Loader implemented in AWS Lambda

● JavaScript ★ 297 🍴 69 Updated 9 days ago

amazon-redshift-udfs

A collection of example UDFs for Amazon Redshift.

● PLpgSQL ★ 86 🍴 18 Updated a day ago

amazon-redshift-monitoring

Amazon Redshift Advanced Monitoring

● Python ★ 63 🍴 19 Updated on Mar 1

References

<https://aws.amazon.com/redshift/developer-resources/>

<https://aws.amazon.com/blogs/aws/quickly-filter-data-in-amazon-redshift-using-interleaved-sorting/>

<http://www.cs.umb.edu/~poneil/StarSchemaB.PDF>

<https://www.slideshare.net/AmazonWebServices/amazon-redshift-optimizing-performance-20150721>

<https://www.quora.com/What-is-the-difference-between-redshift-and-RDS>